

Final Project

December 13, 2021

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from matplotlib.pyplot import figure
from scipy.stats import pearsonr
import numpy.random as rnd
from scipy import stats
```

1 DATA 102 Final Project Fall 2021

Group members: Alisha Mirapuri, Coco Sun, Nameera Faisal Akhtar, Riya Berry

Question 1: (Causal Inference): Does living in states with higher air pollution levels cause an increase in asthma mortality rates?

Question 2: (Comparing GLMs and nonparametric methods): How well does race predict risk for asthma mortality?

1.1 EDA

```
[2]: asthma = pd.read_csv("U.S._Chronic_Disease_Indicators__Asthma.csv")
pollution = pd.read_csv("Daily_Census_Tract-Level_PM2.
↪5_Concentrations__2011-2014.csv")
```

1.1.1 EDA for asthma mortality rates by state

```
[3]: # Exploring what kind of questions exist in our asthma dataset, and how we can
↪use these questions to get asthma mortality rate.
asthma['Question'].unique()
```

```
[3]: array(['Asthma mortality rate',
        'Emergency department visit rate for asthma',
        'Hospitalizations for asthma',
        'Current asthma prevalence among adults aged >= 18 years',
        'Asthma prevalence among women aged 18-44 years',
        'Influenza vaccination among noninstitutionalized adults aged 18-64 years
with asthma',
        'Influenza vaccination among noninstitutionalized adults aged >= 65 years
```

```

with asthma',
    'Pneumococcal vaccination among noninstitutionalized adults aged 18-64
years with asthma',
    'Pneumococcal vaccination among noninstitutionalized adults aged >= 65
years with asthma'],
    dtype=object)

```

```

[4]: # Only looking at those data entries which answer the question of asthma
      ↪ mortality rate, and those that actually have a value.
      # Only looking at the age-adjusted rate as this removes the confounding effect
      ↪ of the age variable.
location = asthma[asthma['Question'] == "Asthma mortality rate"]
location = location[~location['DataValue'].isna()]
location = location[location['DataValueType'] == "Age-adjusted Rate"]
location

```

```

[4]:

```

	YearStart	YearEnd	LocationAbbr	LocationDesc	DataSource	Topic	\
46	2017	2017	US	United States	NVSS	Asthma	
56	2014	2014	CA	California	NVSS	Asthma	
76	2017	2017	CT	Connecticut	NVSS	Asthma	
79	2013	2013	KS	Kansas	NVSS	Asthma	
80	2016	2016	AL	Alabama	NVSS	Asthma	
...	
9772	2013	2013	MI	Michigan	NVSS	Asthma	
9773	2013	2013	NY	New York	NVSS	Asthma	
9775	2013	2013	NC	North Carolina	NVSS	Asthma	
9798	2017	2017	NC	North Carolina	NVSS	Asthma	
9800	2014	2014	TN	Tennessee	NVSS	Asthma	

	Question	Response	DataValueUnit	DataValueType	\
46	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
56	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
76	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
79	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
80	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
...	
9772	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
9773	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
9775	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
9798	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	
9800	Asthma mortality rate	NaN	cases per 1,000,000	Age-adjusted Rate	

	LocationID	TopicID	QuestionID	DataValueTypeID	\
46	...	59	AST	AST4_1	AGEADJRATE
56	...	6	AST	AST4_1	AGEADJRATE
76	...	9	AST	AST4_1	AGEADJRATE
79	...	20	AST	AST4_1	AGEADJRATE

80	...	1	AST	AST4_1	AGEADJRATE
...
9772	...	26	AST	AST4_1	AGEADJRATE
9773	...	36	AST	AST4_1	AGEADJRATE
9775	...	37	AST	AST4_1	AGEADJRATE
9798	...	37	AST	AST4_1	AGEADJRATE
9800	...	47	AST	AST4_1	AGEADJRATE

	StratificationCategoryID1	StratificationID1	StratificationCategoryID2	\
46	OVERALL	OVR		NaN
56	RACE	API		NaN
76	RACE	WHT		NaN
79	RACE	WHT		NaN
80	RACE	WHT		NaN
...	
9772	RACE	WHT		NaN
9773	RACE	BLK		NaN
9775	RACE	WHT		NaN
9798	RACE	WHT		NaN
9800	GENDER	GENF		NaN

	StratificationID2	StratificationCategoryID3	StratificationID3
46	NaN	NaN	NaN
56	NaN	NaN	NaN
76	NaN	NaN	NaN
79	NaN	NaN	NaN
80	NaN	NaN	NaN
...
9772	NaN	NaN	NaN
9773	NaN	NaN	NaN
9775	NaN	NaN	NaN
9798	NaN	NaN	NaN
9800	NaN	NaN	NaN

[1244 rows x 33 columns]

```
[5]: # Checking if location contains the states, as we think it should.
location['LocationDesc'].value_counts()
```

```
[5]: California      56
New York            48
Florida             48
Texas               48
New Jersey          41
Ohio                40
Michigan             40
Illinois            40
```

Georgia	40
Virginia	40
North Carolina	40
Maryland	40
Pennsylvania	40
Tennessee	39
South Carolina	38
Alabama	36
Missouri	35
Arizona	34
Mississippi	34
Indiana	33
Wisconsin	32
Washington	32
Minnesota	30
Massachusetts	30
Louisiana	29
Oregon	27
Oklahoma	27
Colorado	26
Iowa	24
Connecticut	23
Kentucky	22
Utah	21
Nebraska	20
Arkansas	19
Kansas	18
Idaho	12
West Virginia	12
Nevada	9
United States	8
New Mexico	8
Hawaii	4
South Dakota	1

Name: LocationDesc, dtype: int64

```
[6]: # One of the locations is "United States", which we want to get rid of since
      ↪we're looking at mortality rates by state.
```

```
location = location[location['LocationDesc'] != "United States"]
```

```
[7]: # Look at average age-adjusted asthma mortality rates by state.
```

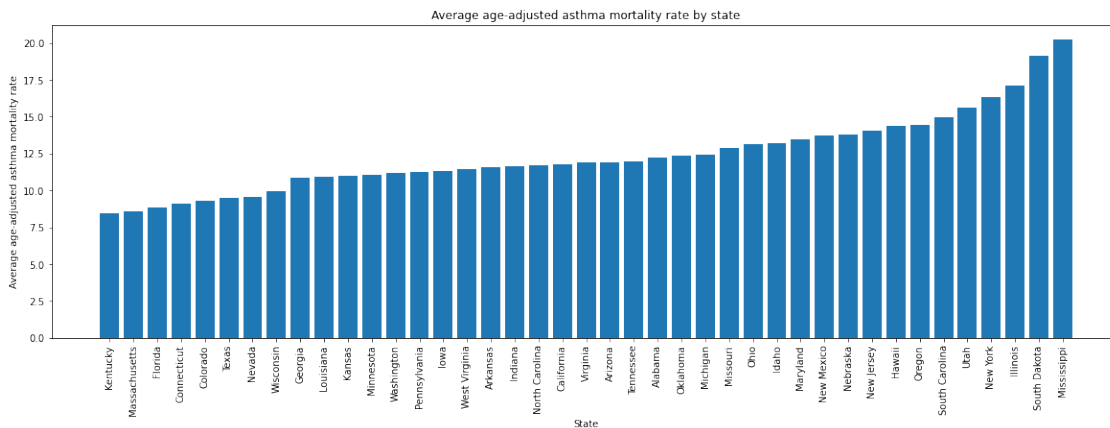
```
location = location[['LocationDesc', 'DataValue']]
```

```
location = location.groupby('LocationDesc').mean()
```

```
location = location.reset_index()
```

```
location = location.sort_values(by='DataValue')
```

```
[8]: # Plotting
plt.figure(figsize=(20, 6))
plt.bar(location['LocationDesc'], location['DataValue'])
plt.xticks(location['LocationDesc'], rotation='vertical');
plt.xlabel("State");
plt.ylabel("Average age-adjusted asthma mortality rate");
plt.title("Average age-adjusted asthma mortality rate by state");
```




Written analysis included in report.

1.1.2 EDA for asthma mortality rates by race

```
[9]: # Only looking at those data entries which answer the question of asthma
      ↳ mortality rate, and those that actually have a value.
      # Only looking at the age-adjusted rate as this removes the confounding effect
      ↳ of the age variable.
      # Only looking at the stratification category of race.
      # Grouping by race, getting the mean of each race.
race = asthma[asthma['StratificationCategoryID1'] == 'RACE']
race = race[race['Question'] == "Asthma mortality rate"]
race = race[race['DataValueType'] == "Age-adjusted Rate"]
race = race[['StratificationID1', 'DataValue']]
race = race[~race['DataValue'].isna()]
race = race.groupby('StratificationID1').mean()
race = race.reset_index()
```

```
[10]: # Plotting
plt.bar(race['StratificationID1'], race['DataValue'])
plt.ylabel("Average age-adjusted asthma mortality rate") ;
plt.xlabel("Race");
plt.title("Average age-adjusted asthma mortality rate by race");
```

```
plt.xticks(race['StratificationID1'], ["API", "Black", "Hispanic", "White"],
           rotation='horizontal')
plt.show();
```




Written analysis included in report.

1.1.3 EDA for asthma hospitalization rates versus race

```
[11]: # Only looking at those data entries which answer the question of asthma
      ↪ hospitalizations, and those that actually have a value.
      # Only looking at the age-adjusted rate as this removes the confounding effect
      ↪ of the age variable.
      # Only looking at the stratification category of race.
      # Grouping by race, getting the mean of each race.
      race = asthma[asthma['StratificationCategoryID1'] == 'RACE']
      race = race[race['Question'] == "Hospitalizations for asthma"]
      race = race[race['DataValueType'] == "Age-adjusted Rate"]
      race = race[['StratificationID1', 'DataValue']]
      race = race[~race['DataValue'].isna()]
      race = race.groupby('StratificationID1').mean()
      race = race.reset_index()
      race
```

```
[11]: StratificationID1  DataValue
0          AIAN      9.614932
1          API       4.445682
2          BLK      19.566720
3          HIS       8.008559
4          WHT       5.806618
```

```
[12]: # Plotting
plt.bar(race['StratificationID1'], race['DataValue'])
plt.ylabel("Average age-adjusted asthma hospitalization rate") ;
plt.xlabel("Race");
plt.title("Average age-adjusted asthma hospitalization rates by race");
plt.xticks(race['StratificationID1'], ["AIAN", "API", "Black", "Hispanic", "White"], rotation='horizontal')
plt.show();
```



1.1.4 EDA for asthma mortality rates versus pollution levels

```
[13]: pollution_by_state = pollution[['statefips', 'ds_pm_pred']]
```

```
[14]: # Looking at mean estimated 24-hour average PM2.5 concentration by state.
pollution_by_state = pollution_by_state.groupby('statefips').mean()
pollution_by_state = pollution_by_state.reset_index()
```

```

pollution_by_state['statefips'] = pollution_by_state['statefips'].astype(int)
pollution_by_state = pollution_by_state[~pollution_by_state['statefips'].isna()]
pollution_by_state

```

```

[14]:
statefips  ds_pm_pred
0          1    12.453149
1          4     6.576452
2          5    11.405171
3          6    11.179666
4          8     6.868177
5          9     7.644805
6         10     9.018218
7         11    10.604991
8         12     7.742284
9         13    11.365693
10        16     7.765538
11        17    12.388659
12        18    12.646612
13        19     8.631860
14        20     9.460055
15        21    12.420450
16        22    10.326207
17        23     6.988289
18        24    10.163588
19        25     7.710368
20        26    10.408863
21        27     6.363210
22        28    11.199066
23        29    10.729569
24        30     7.227646
25        31     8.496382
26        32     7.528881
27        33     7.621899
28        34     9.117607
29        35     6.711759
30        36     9.228158
31        37    10.582490
32        38     5.601993
33        39    12.484810
34        40     9.987869
35        41     6.273876
36        42    10.844855
37        44     6.387972
38        45    10.221010
39        46     6.117873
40        47    11.232846
41        48    11.861321

```


42	49	7.174046
43	50	6.714397
44	51	9.828112
45	53	6.517581
46	54	10.888392
47	55	8.590241
48	56	5.408987

```
[15]: # Looking only at mean, non-null, age-adjusted, asthma mortality rates by state.
asthma_by_state = asthma[asthma['Question'] == "Asthma mortality rate"]
asthma_by_state = asthma_by_state[~asthma_by_state['DataValue'].isna()]
asthma_by_state = asthma_by_state[asthma_by_state['DataValueType'] == "Age-adjusted Rate"]
asthma_by_state = asthma_by_state[asthma_by_state['LocationDesc'] != "United States"]
asthma_by_state = asthma_by_state[['LocationID', 'DataValue']]
asthma_by_state = asthma_by_state.groupby('LocationID').mean()
asthma_by_state = asthma_by_state.reset_index()
asthma_by_state
```

```
[15]:
```


	LocationID	DataValue
0	1	12.236111
1	4	11.923529
2	5	11.605263
3	6	11.746429
4	8	9.307692
5	9	9.100000
6	12	8.833333
7	13	10.880000
8	15	14.400000
9	16	13.200000
10	17	17.125000
11	18	11.612121
12	19	11.287500
13	20	11.022222
14	21	8.450000
15	22	10.941379
16	24	13.495000
17	25	8.590000
18	26	12.412500
19	27	11.036667
20	28	20.226471
21	29	12.891429
22	31	13.805000
23	32	9.566667
24	34	14.019512
25	35	13.750000

26	36	16.350000
27	37	11.720000
28	39	13.172500
29	40	12.370370
30	41	14.418519
31	42	11.255000
32	45	14.939474
33	46	19.100000
34	47	11.943590
35	48	9.466667
36	49	15.590476
37	51	11.880000
38	53	11.203125
39	54	11.450000
40	55	9.943750

```
[16]: merged = pd.merge(asthma_by_state, pollution_by_state, how = 'inner', left_on = 'LocationID', right_on = 'statefips')
```

```
[17]: plt.scatter(merged['ds_pm_pred'], merged['DataValue'])
plt.xlabel("Pollution Rates")
ax = plt.ylabel("Asthma Mortality Rates")
plt.title("Average pollution rates versus average asthma mortality rates for each state");
```

Average pollution rates versus average asthma mortality rates for each state



```
[18]: np.corrcoef(merged['ds_pm_pred'], merged['DataValue'])
```

```
[18]: array([[1.          , 0.00145923],
           [0.00145923, 1.          ]])
```

Written analysis included in report.

1.2 Question 1 (Causal Inference): Does living in states with higher air pollution levels cause an increase in asthma mortality rates?

```
[19]: # Looking at mean estimated 24-hour average PM2.5 concentration by state.
pollution_by_state = pollution[['statefips', 'ds_pm_pred']]
pollution_by_state = pollution_by_state.groupby('statefips').mean()
pollution_by_state = pollution_by_state.reset_index()
pollution_by_state['statefips'] = pollution_by_state['statefips'].astype(int)
pollution_by_state
```

```
[19]:
```

	statefips	ds_pm_pred
0	1	12.453149
1	4	6.576452
2	5	11.405171
3	6	11.179666
4	8	6.868177
5	9	7.644805
6	10	9.018218
7	11	10.604991
8	12	7.742284
9	13	11.365693
10	16	7.765538
11	17	12.388659
12	18	12.646612
13	19	8.631860
14	20	9.460055
15	21	12.420450
16	22	10.326207
17	23	6.988289
18	24	10.163588
19	25	7.710368
20	26	10.408863
21	27	6.363210
22	28	11.199066
23	29	10.729569
24	30	7.227646
25	31	8.496382
26	32	7.528881
27	33	7.621899
28	34	9.117607
29	35	6.711759
30	36	9.228158

31	37	10.582490
32	38	5.601993
33	39	12.484810
34	40	9.987869
35	41	6.273876
36	42	10.844855
37	44	6.387972
38	45	10.221010
39	46	6.117873
40	47	11.232846
41	48	11.861321
42	49	7.174046
43	50	6.714397
44	51	9.828112
45	53	6.517581
46	54	10.888392
47	55	8.590241
48	56	5.408987

```
[20]: # For each state, getting their corresponding mean asthma mortality rates,
# and mean PM2.5 concentration.
merged = pd.merge(asthma_by_state, pollution_by_state, how = 'inner', left_on = 'LocationID', right_on = 'statefips')
```

```
[21]: # This is all states with their corresponding PM2.5 rates, and asthma mortality rates.
merged_all = merged
merged_all
```

```
[21]:
```

	LocationID	DataValue	statefips	ds_pm_pred
0	1	12.236111	1	12.453149
1	4	11.923529	4	6.576452
2	5	11.605263	5	11.405171
3	6	11.746429	6	11.179666
4	8	9.307692	8	6.868177
5	9	9.100000	9	7.644805
6	12	8.833333	12	7.742284
7	13	10.880000	13	11.365693
8	16	13.200000	16	7.765538
9	17	17.125000	17	12.388659
10	18	11.612121	18	12.646612
11	19	11.287500	19	8.631860
12	20	11.022222	20	9.460055
13	21	8.450000	21	12.420450
14	22	10.941379	22	10.326207
15	24	13.495000	24	10.163588
16	25	8.590000	25	7.710368

17	26	12.412500	26	10.408863
18	27	11.036667	27	6.363210
19	28	20.226471	28	11.199066
20	29	12.891429	29	10.729569
21	31	13.805000	31	8.496382
22	32	9.566667	32	7.528881
23	34	14.019512	34	9.117607
24	35	13.750000	35	6.711759
25	36	16.350000	36	9.228158
26	37	11.720000	37	10.582490
27	39	13.172500	39	12.484810
28	40	12.370370	40	9.987869
29	41	14.418519	41	6.273876
30	42	11.255000	42	10.844855
31	45	14.939474	45	10.221010
32	46	19.100000	46	6.117873
33	47	11.943590	47	11.232846
34	48	9.466667	48	11.861321
35	49	15.590476	49	7.174046
36	51	11.880000	51	9.828112
37	53	11.203125	53	6.517581
38	54	11.450000	54	10.888392
39	55	9.943750	55	8.590241

```
[22]: # This is the most industrialized states (California, New Jersey, Texas, New
      ↪York, Florida) with their corresponding PM2.5 rates,
      # and asthma mortality rates.
merged_some = merged[merged['LocationID'].isin([6, 12, 48, 36, 34])]
merged_some
```

```
[22]:   LocationID  DataValue  statefips  ds_pm_pred
3           6   11.746429          6   11.179666
6          12    8.833333          12    7.742284
23         34   14.019512          34    9.117607
25         36   16.350000          36    9.228158
34         48    9.466667          48   11.861321
```

In this part, we will set up our causal inference problem as follows.

The state's average age-adjusted asthma mortality rate score is linear in the state's average PM2.5 concentration levels and the state's level of industrialization.

$$Y = \beta_1 Z + \beta_2 X + \epsilon$$

where,

Z = the state's average PM2.5 concentration levels, and

Y = the state's average asthma mortality rate.

The degree of industrialization/population of a state X affects both Z and Y , but is not observed. As a result, we will estimate the causal effect by using plain linear regression (OLS) on the observed variables and Y . We will also use an intercept term.

As a result, the equation we're solving for becomes:

$$\hat{\beta}_1, \hat{c}_1 = \arg \min_{\beta_1, c_1} \|Y - \beta_1 Z - c_1\|_2^2$$

```
[23]: # Using code from lab 7.

import statsmodels.api as sm

def fit_OLS_model(df, target_variable, explanatory_variables, intercept =
↪False):
    """
    Fits an OLS model from data.

    Inputs:
        df: pandas DataFrame
        target_variable: string, name of the target variable
        explanatory_variables: list of strings, names of the explanatory
↪variables
        intercept: bool, if True add intercept term

    Outputs:
        fitted_model: model containing OLS regression results
    """

    target = df[target_variable]
    inputs = df[explanatory_variables]
    if intercept:
        inputs = sm.add_constant(inputs)

    fitted_model = sm.OLS(target, inputs).fit()
    return(fitted_model)
```

```
[24]: # Fitting the model for all states
gammas_model_all = fit_OLS_model(merged_all, 'DataValue', 'ds_pm_pred',
↪intercept=True);
print(gammas_model_all.summary());
```

OLS Regression Results			
=====			
Dep. Variable:	DataValue	R-squared:	0.000
Model:	OLS	Adj. R-squared:	-0.026
Method:	Least Squares	F-statistic:	8.092e-05
Date:	Mon, 13 Dec 2021	Prob (F-statistic):	0.993
Time:	19:03:25	Log-Likelihood:	-95.364

```

No. Observations:      40    AIC:      194.7
Df Residuals:          38    BIC:      198.1
Df Model:              1
Covariance Type:      nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	12.3287	2.040	6.043	0.000	8.198	16.459
ds_pm_pred	0.0019	0.211	0.009	0.993	-0.424	0.428
Omnibus:	10.167		Durbin-Watson:	1.764		
Prob(Omnibus):	0.006		Jarque-Bera (JB):	9.348		
Skew:	1.038		Prob(JB):	0.00934		
Kurtosis:	4.141		Cond. No.	46.9		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

/opt/conda/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142:

FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
x = pd.concat(x[:, :order], 1)
```

In the last part, we talked about how the degree of industrialization/population of a state X affects both Z and Y , but is not observed. In this part, we will only use those states that have a comparable/similar degree of industrialization, in order to minimize this confounding effect.

Let $\hat{\beta}_s$ and \hat{c}_s be the parameters for this new model where the subscript s denotes the fact that only **some** states are used in this model.

Here, the equation we're solving for becomes:

$$\hat{\beta}_s, \hat{c}_s = \arg \min_{\beta_s, c_s} \|Y - \beta_s Z - c_s\|_2^2$$

where,

Z = the state's average PM2.5 concentration levels.

Y = the state's average asthma mortality rate

like before.

```

[25]: # Fitting the model for some (most industrialized) states
gammas_model_some = fit_OLS_model(merged_some, 'DataValue', 'ds_pm_pred',
    ↪ intercept=True)
print(gammas_model_some.summary())

```

OLS Regression Results

```

=====
Dep. Variable:      DataValue    R-squared:      0.019

```

```

Model:                OLS      Adj. R-squared:      -0.307
Method:               Least Squares    F-statistic:      0.05962
Date:                 Mon, 13 Dec 2021    Prob (F-statistic):    0.823
Time:                 19:03:25    Log-Likelihood:      -12.211
No. Observations:      5    AIC:                28.42
Df Residuals:          3    BIC:                27.64
Df Model:              1
Covariance Type:       nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          14.6610      10.679        1.373      0.263     -19.324      48.646
ds_pm_pred      -0.2624       1.074       -0.244      0.823      -3.682       3.157
=====
Omnibus:                nan    Durbin-Watson:                2.305
Prob(Omnibus):           nan    Jarque-Bera (JB):                0.340
Skew:                    0.103    Prob(JB):                0.844
Kurtosis:                1.740    Cond. No.                66.7
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

/opt/conda/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
    x = pd.concat(x[:, :order], 1)
/opt/conda/lib/python3.9/site-packages/statsmodels/stats/stattools.py:74:
ValueWarning: omni_normtest is not valid with less than 8 observations; 5
samples were given.
    warn("omni_normtest is not valid with less than 8 observations; %i "

```

```
[26]: print(gammas_model_all.params[1], gammas_model_all.params[0])
```

```
0.0018936472133398718 12.328733568284406
```

```
[27]: print(gammas_model_some.params[1], gammas_model_some.params[0])
```

```
-0.2623560729123282 14.66104830450967
```

As a result, we have:

$$\hat{\beta}_1, \hat{c}_1 = 0.0018936472133398718, 12.328733568284406$$

$$\hat{\beta}_s, \hat{c}_s = -0.2623560729123282, 14.66104830450967$$

```
[28]: gammas_model_all.bse
```



```
[28]: const          2.040299
      ds_pm_pred     0.210516
      dtype: float64
```

```
[29]: gammas_model_some.bse
```

```
[29]: const          10.678931
      ds_pm_pred     1.074458
      dtype: float64
```

Written analysis included in report.

1.3 Question 2 (Comparing GLMs and nonparametric methods): How well does race predict risk for asthma mortality?

```
[30]: race = asthma[asthma['StratificationCategoryID1'] == 'RACE']
      race = race[race['Question'] == "Asthma mortality rate"]
      race = race[~race['DataValue'].isna()]
      race = race[race['DataValueType'] == 'Age-adjusted Rate']
      race = race[['StratificationID1', 'DataValue']]
      race = pd.get_dummies(race)
      race
```

```
[30]:
```

	DataValue	StratificationID1_API	StratificationID1_BLK	\
56	9.7	1	0	
76	5.7	0	0	
79	10.3	0	0	
80	6.7	0	0	
112	7.1	0	0	
...	
9744	19.2	0	0	
9772	7.1	0	0	
9773	34.4	0	1	
9775	7.8	0	0	
9798	8.3	0	0	

	StratificationID1_HIS	StratificationID1_WHT
56	0	0
76	0	1
79	0	1
80	0	1
112	0	1
...
9744	1	0
9772	0	1
9773	0	0
9775	0	1

9798

0

1

[466 rows x 5 columns]

1.3.1 Non-parametric method for Research Question 2

Here, we decided to use a decision tree model to predict risk of asthma mortality based on race.

```
[31]: # Performing the train-test split.
from sklearn.model_selection import train_test_split

train, test = train_test_split(race, test_size=0.30, random_state=102)
X_train = train.iloc[:, 1:]
y_train = train['DataValue']

[32]: # Fitting the model and using it to predict.
from sklearn.ensemble import RandomForestRegressor

random_forest_model = RandomForestRegressor(max_features=1).fit(X_train,
    ↪y_train)

train["forest_pred"] = random_forest_model.predict(X_train)
test["forest_pred"] = random_forest_model.predict(test.iloc[:, 1:])

/tmp/ipykernel_392/698277838.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    train["forest_pred"] = random_forest_model.predict(X_train)
/tmp/ipykernel_392/698277838.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    test["forest_pred"] = random_forest_model.predict(test.iloc[:, 1:])

[33]: # Evaluating the model.
train_rmse = np.mean((train["forest_pred"] - train["DataValue"]) ** 2) ** 0.5
test_rmse = np.mean((test["forest_pred"] - test["DataValue"]) ** 2) ** 0.5

print("Training set error for random forest:", train_rmse)
print("Test set error for random forest:    ", test_rmse)
```

Training set error for random forest: 5.06019303437648

Test set error for random forest: 4.601109416245368

```
[34]: # Evaluating the model on the training set
random_forest_model.score(X_train, y_train)
```

```
[34]: 0.6703969849829345
```

```
[35]: # Evaluating the model on the test set
X_test = test.iloc[:, 1:5]
y_test = test['DataValue']
random_forest_model.score(X_test, y_test)
```

```
[35]: 0.746908930274781
```

```
[36]: pd.value_counts(test['forest_pred'])
```

```
[36]: 9.128748      82
      24.951027   43
      11.456613   12
      9.645222    3
      Name: forest_pred, dtype: int64
```

1.3.2 GLM for Research Question 2

```
[37]: # Fitting the GLM, and displaying summary.
import statsmodels.api as sm
gaussian_model = sm.GLM(
    train.DataValue, sm.add_constant(train.iloc[:, 1:5]),
    family=sm.families.Gaussian()
)
gaussian_results = gaussian_model.fit()
print(gaussian_results.summary())
```

```

                        Generalized Linear Model Regression Results
=====
Dep. Variable:          DataValue   No. Observations:          326
Model:                  GLM        Df Residuals:                322
Model Family:           Gaussian   Df Model:                  3
Link Function:          identity    Scale:                    25.932
Method:                 IRLS       Log-Likelihood:           -991.12
Date:                   Mon, 13 Dec 2021   Deviance:                 8345.7
Time:                   19:03:26    Pearson chi2:             8.35e+03
No. Iterations:         3
Covariance Type:        nonrobust
=====
=====
                        coef      std err          z      P>|z|      [0.025
0.975]
-----
```

```

-----
const          11.0143      0.518      21.281      0.000      10.000
12.029
StratificationID1_API  -1.4143      1.838      -0.769      0.442      -5.017
2.189
StratificationID1_BLK   13.8126      0.646      21.378      0.000      12.546
15.079
StratificationID1_HIS    0.4901      0.972      0.504      0.614      -1.415
2.395
StratificationID1_WHT   -1.8741      0.590      -3.177      0.001      -3.030
-0.718
=====
=====

```

```

/opt/conda/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
    x = pd.concat(x[:, :order], 1)

```

We will use bootstrapped Confidence Intervals to estimate uncertainty.

```

[38]: # For evaluating our model, we are using code from Lecture 12.
def bootstrap_xy(X, y, fnc, w=None, B=1000, plot=True):
    d = X.shape[1]
    N = X.shape[0]
    w_hat = fnc(X, y)
    w_boot = np.zeros(shape=(B,d))
    for b in range(B):
        bootstrap_indices = rnd.choice(np.arange(N), N)
        bootstrap_X = X.iloc[bootstrap_indices, :]
        bootstrap_y = y.iloc[bootstrap_indices]
        w_boot[b,:] = fnc(bootstrap_X, bootstrap_y)
    if plot:
        plt.scatter(w_boot[:,0], w_boot[:,1], c='b')
        plt.scatter(w_hat[0], w_hat[1], c='r', marker='x', s=300)
        if w:
            plt.scatter(w[0], w[1], c='g', marker='x', s=300)
        plt.show()
    return w_boot


def lin_model(x, y):
    model = sm.GLM(
        y, x,
        family=sm.families.Gaussian()
    )
    results = model.fit()
    params = results.params

```

```
return params
```

```
[39]: w_gaussian_boot = bootstrap_xy(sm.add_constant(train.iloc[:, 1:5]), train.  
    ↪DataValue, lin_model);
```

```
/opt/conda/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142:  
FutureWarning: In a future version of pandas all arguments of concat except for  
the argument 'objs' will be keyword-only  
x = pd.concat(x[:, :order], 1)
```



```
[40]: beta_0, beta_1, beta_2, beta_3, beta_4 = w_gaussian_boot.std(axis = 0)  
print(f"Bootstrap std error for constant: {beta_0:.3f}")  
print(f"Bootstrap std error for API individuals: {beta_1:.3f}")  
print(f"Bootstrap std error for Black individuals: {beta_2:.3f}")  
print(f"Bootstrap std error for Hispanic individuals: {beta_3:.3f}")  
print(f"Bootstrap std error for White individuals: {beta_4:.3f}")
```

```
Bootstrap std error for constant: 0.343  
Bootstrap std error for API individuals: 0.598  
Bootstrap std error for Black individuals: 0.672  
Bootstrap std error for Hispanic individuals: 1.166  
Bootstrap std error for White individuals: 0.373
```

```
[41]: gaussian_table = pd.DataFrame(w_gaussian_boot, columns = ["Constant", "API", "  
    ↪Black", "Hispanic", "White"])  
gaussian_table
```

```
[41]:
```

	Constant	API	Black	Hispanic	White
0	10.704273	-0.654273	14.407545	-1.109829	-1.939169
1	11.363480	-1.396813	14.604288	0.100520	-1.944515
2	11.163439	-1.996773	13.598508	1.531561	-1.969856
3	10.980597	-2.020597	12.391366	2.571577	-1.961749
4	10.620586	-1.360586	13.911453	-0.331697	-1.598586
..
995	11.499239	-2.124239	13.009457	3.263261	-2.649239
996	11.071893	-0.871893	13.372551	0.275726	-1.704490
997	11.556859	-1.381859	13.281006	2.181237	-2.523525
998	11.020961	-1.770961	14.096933	0.683584	-1.988594
999	10.918679	-1.352012	14.967528	-1.023441	-1.673396

[1000 rows x 5 columns]

```
[42]: # 95% confidence interval for the constant
stats.norm.interval(0.95, loc=np.mean(gaussian_table["Constant"]), scale= np.
↳std(gaussian_table["Constant"]))
```

```
[42]: (10.346926074409838, 11.690320691160823)
```

```
[43]: # 95% confidence interval for the API coefficient
stats.norm.interval(0.95, loc=np.mean(gaussian_table["API"]), scale= np.
↳std(gaussian_table["API"]))
```

```
[43]: (-2.5721384935880436, -0.22828408640858822)
```

```
[44]: # 95% confidence interval for the Black coefficient
stats.norm.interval(0.95, loc=np.mean(gaussian_table["Black"]), scale= np.
↳std(gaussian_table["Black"]))
```

```
[44]: (12.503099780224895, 15.138469855950566)
```

```
[45]: # 95% confidence interval for the Hispanic coefficient
stats.norm.interval(0.95, loc=np.mean(gaussian_table["Hispanic"]), scale= np.
↳std(gaussian_table["Hispanic"]))
```

```
[45]: (-1.7982929912615835, 2.7723986603399866)
```

```
[46]: # 95% confidence interval for the White coefficient
stats.norm.interval(0.95, loc=np.mean(gaussian_table["White"]), scale= np.
↳std(gaussian_table["White"]))
```

```
[46]: (-2.6201746704856292, -1.1578312892009555)
```